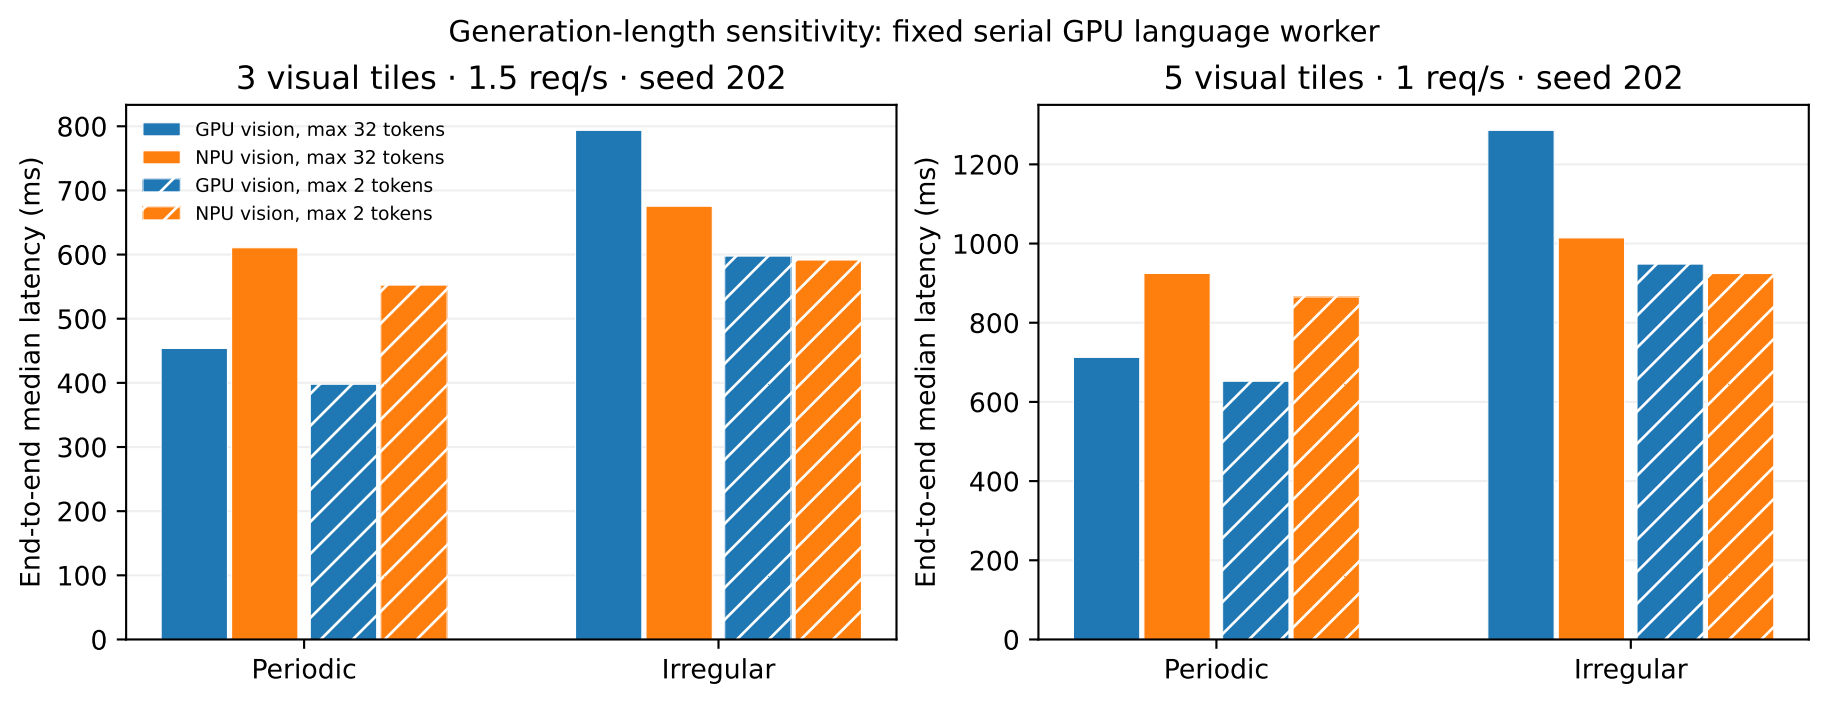
Generation-length sensitivity: fixed serial GPU language worker
3 visual tiles · 1.5 req/s · seed 202
5 visual tiles · 1 req/s · seed 202
End-to-end median latency (ms)
GPU vision, max 32 tokens
NPU vision, max 32 tokens
GPU vision, max 2 tokens
NPU vision, max 2 tokens
0
100
200
300
400
500
600
700
800
1000
1200
Periodic
Irregular